

01-oppgave-beskrivelse.md

Operating Systems - Assignment 1

Task 1:

1. Describe user-mode and kernel-mode, and why.

A mode switch from kernel to user-mode must happen because user applications are not allowed to execute certain CPU instructions or access certain areas of memory directly. This is a security and stability measure to prevent user applications from interfering with the core functionality of the operating system.

2. Program "printx.c"

Manual References

Explanation: Here you use `man 3 <function>` to find information about the functions you use in your code. The number 3 refers to the C library manual section.

Manual References

- `man 3 printf`

DESCRIPTION

The functions in the `printf()` family produce output according to a format as described below. The functions `printf()` and `vprintf()` write output to `stdout`, the standard

output stream; `fprintf()` and `vfprintf()` write output to the given output stream; `sprintf()`, `snprintf()`, `vsprintf()`, and `vsnprintf()` write to the character string `str`.

The function `dprintf()` is the same as `fprintf()` except that it outputs to a file descriptor, `fd`, instead of to a `stdio(3)` stream.

The functions `snprintf()` and `vsnprintf()` write at most `size` bytes (including the terminating null byte `'\0'`) to `str`.

The functions `vprintf()`, `vfprintf()`, `vdprintf()`, `vsprintf()`, `vsnprintf()` are equivalent to the functions `printf()`, `fprintf()`, `dprintf()`, `sprintf()`, `snprintf()`, respec-

tively, except that they are called with a `va_list` instead of a variable number of arguments. These functions do not call the `va_end` macro. Because they invoke the

`va_arg` macro, the value of `ap` is undefined after the call. See `stdarg(3)`.

All of these functions write the output under the control of a format string that specifies how subsequent arguments (or arguments accessed via the variable-length argu-

ment facilities of `stdarg(3)`) are converted for output.

C99 and POSIX.1-2001 specify that the results are undefined if a call to `sprintf()`, `snprintf()`, `vsprintf()`, or `vsnprintf()` would cause copying to take place between ob-

jects that overlap (e.g., if the target string array and one of the supplied input arguments refer to the same buffer). See CAVEATS.

- **man 3 fgets**

DESCRIPTION

`fgetc()` reads the next character from stream and returns it as an unsigned char cast to an int, or EOF on end of file or error.

`getc()` is equivalent to `fgetc()` except that it may be implemented as a macro which evaluates stream more than once.

`getchar()` is equivalent to `getc(stdin)`.

`fgets()` reads in at most one less than size characters from stream and stores them into the buffer pointed to by s. Reading stops after an EOF or a newline. If a new-

line is read, it is stored into the buffer. A terminating null byte ('\0') is stored after the last character in the buffer.

`ungetc()` pushes c back to stream, cast to unsigned char, where it is available for subsequent read operations. Pushed-back characters will be returned in reverse order;

only one pushback is guaranteed.

Calls to the functions described here can be mixed with each other and with calls to other input functions from the stdio library for the same input stream.

For nonlocking counterparts, see `unlocked_stdio(3)`.

- **man 3 atoi**

DESCRIPTION

`fgetc()` reads the next character from stream and returns it as an unsigned char cast to an int, or EOF on end of file or error.

`getc()` is equivalent to `fgetc()` except that it may be implemented as a macro which evaluates stream more than once.

`getchar()` is equivalent to `getc(stdin)`.

`fgets()` reads in at most one less than size characters from stream and stores them into the buffer pointed to by s. Reading stops after an EOF or a newline. If a new-

line is read, it is stored into the buffer. A terminating null byte ('\0') is stored after the last character in the buffer.

`ungetc()` pushes c back to stream, cast to unsigned char, where it is available for subsequent read operations. Pushed-back characters will be returned in reverse order;

only one pushback is guaranteed.

Calls to the functions described here can be mixed with each other and with calls to other input functions from the stdio library for the same input stream.

For nonlocking counterparts, see `unlocked_stdio(3)`.

- **man 3 sscanf**

DESCRIPTION

The `sscanf()` family of functions scans formatted input according to format as described below. This format may contain conversion specifications; the results from such

conversions, if any, are stored in the locations pointed to by the pointer arguments that follow format. Each pointer argument must be of a type that is appropriate for

the value returned by the corresponding conversion specification.

If the number of conversion specifications in format exceeds the number of pointer arguments, the results are undefined. If the number of pointer arguments exceeds the

number of conversion specifications, then the excess pointer arguments are evaluated, but are otherwise ignored.

`sscanf()` These functions read their input from the string pointed to by str.

The `vsscanf()` function is analogous to `vsprintf(3)`.

The format string consists of a sequence of directives which describe how to process the sequence of input characters. If processing of a directive fails, no further in-

put is read, and `sscanf()` returns. A "failure" can be either of the following: input failure, meaning that input characters were unavailable, or matching failure, meaning that the input was inappropriate (see below).

A directive is one of the following:

- A sequence of white-space characters (space, tab, newline, etc.; see `isspace(3)`). This directive matches any amount of white space, including none, in the input.
- An ordinary character (i.e., one other than white space or '%'). This character must exactly match the next character of input.
- A conversion specification, which commences with a '%' (percent) character. A sequence of characters from the input is converted according to this specification, and the result is placed in the corresponding pointer argument. If the next item of input does not match the conversion specification, the conversion fails—this is a matching failure.

Each conversion specification in format begins with either the character '%' or the character sequence "%n\$" (see below for the distinction) followed by:

- An optional '*' assignment-suppression character: `sscanf()` reads input as directed by the conversion specification, but discards the input. No corresponding pointer argument is required, and this specification is not included in the count of successful assignments returned by `scanf()`.
- For decimal conversions, an optional quote character ('). This specifies that the input number may include thousands' separators as defined by the `LC_NUMERIC` category of the current locale. (See `setlocale(3)`.) The quote character may precede or follow the '*' assignment-suppression character.
- An optional 'm' character. This is used with string conversions (`%s`, `%c`, `%[]`), and relieves the caller of the need to allocate a corresponding buffer to hold the input: instead, `sscanf()` allocates a buffer of sufficient

size, and assigns the address of this buffer to the corresponding pointer argument, which should be a pointer to a char * variable (this variable does not need to be initialized before the call). The caller should subsequently free(3) this buffer when it is no longer required.

- An optional decimal integer which specifies the maximum field width. Reading of characters stops either when this maximum is reached or when a nonmatching character is found, whichever happens first. Most conversions discard initial white space characters (the exceptions are noted below), and these discarded characters don't count toward the maximum field width. String input conversions store a terminating null byte ('\0') to mark the end of the input; the maximum field width does not include this terminator.
- An optional type modifier character. For example, the l type modifier is used with integer conversions such as %d to specify that the corresponding pointer argument refers to a long rather than a pointer to an int.
- A conversion specifier that specifies the type of input conversion to be performed.

The conversion specifications in format are of two forms, either beginning with '%' or beginning with "%n\$". The two forms should not be mixed in the same format string, except that a string containing "%n\$" specifications can include %% and %*. If format contains '%' specifications, then these correspond in order with successive pointer arguments. In the "%n\$" form (which is specified in POSIX.1-2001, but not C99), n is a decimal integer that specifies that the converted input should be placed in the location referred to by the n-th pointer argument following format.

Task 2 memory and segments.

1. Describe the different memory segments in my code.

Address: 0x5a3b79218014; Value: 0

Address: 0x7ffd17f5c9ac; Value: 1

Address: 0x7ffd17f5c9b0; Address: 0x5a3ba69012a0; Value: 2

By looking at the print of the program 01-mem.bin we can come to this conclusion:

`var1` is a global variable, in code it is allocated before the main function, and the address is `0x5...` which is a low address, so it is in the data segment.

`var2` is a local variable, it is allocated on the stack when the main function is called, and the address is `0x7...` which is a higher address, so it is in the stack segment.

`var3` is a pointer to a dynamically allocated memory on the heap, the pointer itself is stored on the stack (address `0x7...`), but it points to a memory location on the heap (address `0x5...`), so it is in the heap segment.

2. Describe the purpose of a process' address space, starting with high addresses to low addresses.

Each segment has its own purpose:

- **Stack segment:**, in this case `0x7...`, is used for local variables and function call management. With dynamic memory allocation.
- **Heap segment:**, in this case `0x5...`, is used for dynamic memory allocation outside of the stack, and remains allocated until it is explicitly freed. (e.g., using `malloc` and `free` in C). Example of this in the 01-mem.c code is the line `var3 = malloc(sizeof(int));`. For good practice, we should always free the memory we have allocated to the heap, but using `free(var3);` when we are done using it.
- **Data segment:**, in this case `0x5...`, is used for global and static variables. This segment is divided into initialized and uninitialized sections. Initialized global and static variables are stored in the initialized data segment, while uninitialized ones are stored in the uninitialized data segment (BSS).
- **Text segment:**, not shown in the example, is used for the actual code of the program. This segment is typically read-only to prevent accidental modification of instructions.
- **0x0 - NULL pointer:**, the address `0x0` is reserved to represent a null pointer, which indicates that a pointer does not point to any valid memory location.

3. Differences between static, global and local variables.

Global variables:

- Declared outside of functions. (*including main*)

- Accessible from any function within the same file.
- Stored in the data segment of memory.
- **Can be accessed from other files using the `extern` keyword.**

Static variables:

- Declared with the `static` keyword.
- Can be declared inside or outside of functions.
- Functions like global variables, but their scope is limited to the file if declared outside a function, or to the function if declared inside a function.
- Also stored in the data segment of memory.
- **Cannot be accessed from other files.**
- Declared outside a function it acts much like a global variable, if we exclude `extern`.

Local variables:

- Declared inside a function.
- Only accessible within the function they are declared in.
- Stored on the stack.
- **Cannot be accessed from other functions or files.**
- Gets destroyed when the function exits.
- *Another way of describing it, is that local variables exist only while the function is executing, similar to how variables in `main` exist only during the program's execution. They are created when the function starts and destroyed when it ends. They are isolated and treat the function's stack frame as their own environment.*

Task 3: Program code:

1. Determine the sizes of text, data, bss.

text	data	bss	dec	hex	filename
1799	616	8	2423	977	./01-mem.bin

- **Text segment:** 1799 bytes
- **Data segment:** 616 bytes
- **BSS segment:** 8 bytes

2. Find the adress.

```
01-mem.bin:      file format elf64-x86-64
architecture: i386:x86-64, flags 0x000000150:
HAS_SYMS, DYNAMIC, D_PAGED
start address 0x000000000000010a0
```

Start address: 0x000000000000010a0

3. Dissassemble compiled program & Capture output and find name of the function at the start address.

```
01-mem.bin:      file format elf64-x86-64
```

Disassembly of section .init:

00000000000001000 <_init>:

1000:	f3 0f 1e fa	endbr64	
1004:	48 83 ec 08	sub	\$0x8,%rsp
1008:	48 8b 05 d9 2f 00 00	mov	0x2fd9(%rip),%rax # 3fe
100f:	48 85 c0	test	%rax,%rax
1012:	74 02	je	1016 <_init+0x16>
1014:	ff d0	call	*%rax
1016:	48 83 c4 08	add	\$0x8,%rsp
101a:	c3	ret	

Disassembly of section .plt:

00000000000001020 <.plt>:

1020:	ff 35 8a 2f 00 00	push	0x2f8a(%rip) # 3fb0 <_G
1026:	ff 25 8c 2f 00 00	jmp	*0x2f8c(%rip) # 3fb8 <_G
102c:	0f 1f 40 00	nopl	0x0(%rax)
1030:	f3 0f 1e fa	endbr64	
1034:	68 00 00 00 00	push	\$0x0
1039:	e9 e2 ff ff ff	jmp	1020 <_init+0x20>
103e:	66 90	xchg	%ax,%ax
1040:	f3 0f 1e fa	endbr64	
1044:	68 01 00 00 00	push	\$0x1
1049:	e9 d2 ff ff ff	jmp	1020 <_init+0x20>
104e:	66 90	xchg	%ax,%ax
1050:	f3 0f 1e fa	endbr64	
1054:	68 02 00 00 00	push	\$0x2
1059:	e9 c2 ff ff ff	jmp	1020 <_init+0x20>
105e:	66 90	xchg	%ax,%ax

Disassembly of section .plt.got:

00000000000001060 <__cxa_finalize@plt>:

```

1060:      f3 0f 1e fa      endbr64
1064:      ff 25 8e 2f 00 00    jmp     *0x2f8e(%rip)      # 3ff8 <_
106a:      66 0f 1f 44 00 00    nopw   0x0(%rax,%rax,1)

```

Disassembly of section .plt.sec:

00000000000001070 <__stack_chk_fail@plt>:

```

1070:      f3 0f 1e fa      endbr64
1074:      ff 25 46 2f 00 00    jmp     *0x2f46(%rip)      # 3fc0 <_
107a:      66 0f 1f 44 00 00    nopw   0x0(%rax,%rax,1)

```

00000000000001080 <printf@plt>:

```

1080:      f3 0f 1e fa      endbr64
1084:      ff 25 3e 2f 00 00    jmp     *0x2f3e(%rip)      # 3fc8 <p
108a:      66 0f 1f 44 00 00    nopw   0x0(%rax,%rax,1)

```

00000000000001090 <malloc@plt>:

```

1090:      f3 0f 1e fa      endbr64
1094:      ff 25 36 2f 00 00    jmp     *0x2f36(%rip)      # 3fd0 <m
109a:      66 0f 1f 44 00 00    nopw   0x0(%rax,%rax,1)

```

Disassembly of section .text:

000000000000010a0 <_start>:

```

10a0:      f3 0f 1e fa      endbr64
10a4:      31 ed              xor     %ebp,%ebp
10a6:      49 89 d1              mov     %rdx,%r9
10a9:      5e                    pop     %rsi
10aa:      48 89 e2              mov     %rsp,%rdx
10ad:      48 83 e4 f0           and     $0xfffffffffffffff0,%rsp
10b1:      50                    push    %rax
10b2:      54                    push    %rsp
10b3:      45 31 c0              xor     %r8d,%r8d
10b6:      31 c9              xor     %ecx,%ecx
10b8:      48 8d 3d ca 00 00 00    lea     0xca(%rip),%rdi      # 1189
10bf:      ff 15 13 2f 00 00      call    *0x2f13(%rip)      # 3fd8 <_
10c5:      f4                    hlt
10c6:      66 2e 0f 1f 84 00 00    cs nopw 0x0(%rax,%rax,1)
10cd:      00 00 00

```

000000000000010d0 <deregister_tm_clones>:

```

10d0:      48 8d 3d 39 2f 00 00    lea     0x2f39(%rip),%rdi      # 401
10d7:      48 8d 05 32 2f 00 00    lea     0x2f32(%rip),%rax      # 401
10de:      48 39 f8              cmp     %rdi,%rax
10e1:      74 15              je      10f8 <deregister_tm_clones+0x2
10e3:      48 8b 05 f6 2e 00 00    mov     0x2ef6(%rip),%rax      # 3fe

```

```

10ea:      48 85 c0                test    %rax,%rax
10ed:      74 09                  je      10f8 <deregister_tm_clones+0x2
10ef:      ff e0                  jmp     *%rax
10f1:      0f 1f 80 00 00 00 00    nopl    0x0(%rax)
10f8:      c3                      ret
10f9:      0f 1f 80 00 00 00 00    nopl    0x0(%rax)

00000000000001100 <register_tm_clones>:
1100:      48 8d 3d 09 2f 00 00    lea     0x2f09(%rip),%rdi        # 401
1107:      48 8d 35 02 2f 00 00    lea     0x2f02(%rip),%rsi        # 401
110e:      48 29 fe                  sub     %rdi,%rsi
1111:      48 89 f0                  mov     %rsi,%rax
1114:      48 c1 ee 3f              shr     $0x3f,%rsi
1118:      48 c1 f8 03              sar     $0x3,%rax
111c:      48 01 c6                  add     %rax,%rsi
111f:      48 d1 fe                  sar     $1,%rsi
1122:      74 14                      je      1138 <register_tm_clones+0x38>
1124:      48 8b 05 c5 2e 00 00    mov     0x2ec5(%rip),%rax        # 3ff
112b:      48 85 c0                test    %rax,%rax
112e:      74 08                      je      1138 <register_tm_clones+0x38>
1130:      ff e0                  jmp     *%rax
1132:      66 0f 1f 44 00 00        nopw    0x0(%rax,%rax,1)
1138:      c3                      ret
1139:      0f 1f 80 00 00 00 00    nopl    0x0(%rax)

00000000000001140 <__do_global_dtors_aux>:
1140:      f3 0f 1e fa              endbr64
1144:      80 3d c5 2e 00 00 00    cmpb    $0x0,0x2ec5(%rip)        # 401
114b:      75 2b                      jne     1178 <__do_global_dtors_aux+0x
114d:      55                          push    %rbp
114e:      48 83 3d a2 2e 00 00    cmpq    $0x0,0x2ea2(%rip)        # 3ff
1155:      00
1156:      48 89 e5                  mov     %rsp,%rbp
1159:      74 0c                      je      1167 <__do_global_dtors_aux+0x
115b:      48 8b 3d a6 2e 00 00    mov     0x2ea6(%rip),%rdi        # 400
1162:      e8 f9 fe ff ff          call    1060 <__cxa_finalize@plt>
1167:      e8 64 ff ff ff          call    10d0 <deregister_tm_clones>
116c:      c6 05 9d 2e 00 00 01    movb    $0x1,0x2e9d(%rip)        # 401
1173:      5d                          pop     %rbp
1174:      c3                      ret
1175:      0f 1f 00                  nopl    (%rax)
1178:      c3                      ret
1179:      0f 1f 80 00 00 00 00    nopl    0x0(%rax)

00000000000001180 <frame_dummy>:
1180:      f3 0f 1e fa              endbr64
1184:      e9 77 ff ff ff          jmp     1100 <register_tm_clones>

00000000000001189 <main>:

```

```

1189:      f3 0f 1e fa                endbr64
118d:      55                        push    %rbp
118e:      48 89 e5                  mov     %rsp,%rbp
1191:      48 83 ec 20               sub     $0x20,%rsp
1195:      64 48 8b 04 25 28 00      mov     %fs:0x28,%rax
119c:      00 00
119e:      48 89 45 f8                mov     %rax,-0x8(%rbp)
11a2:      31 c0                      xor     %eax,%eax
11a4:      c7 45 ec 01 00 00 00      movl    $0x1,-0x14(%rbp)
11ab:      bf 04 00 00 00             mov     $0x4,%edi
11b0:      e8 db fe ff ff             call    1090 <malloc@plt>
11b5:      48 89 45 f0                mov     %rax,-0x10(%rbp)
11b9:      48 8b 45 f0                mov     -0x10(%rbp),%rax
11bd:      c7 00 02 00 00 00          movl    $0x2,(%rax)
11c3:      8b 05 4b 2e 00 00          mov     0x2e4b(%rip),%eax      # 401
11c9:      89 c2                      mov     %eax,%edx
11cb:      48 8d 05 42 2e 00 00      lea     0x2e42(%rip),%rax      # 401
11d2:      48 89 c6                  mov     %rax,%rsi
11d5:      48 8d 05 2c 0e 00 00      lea     0xe2c(%rip),%rax      # 2008
11dc:      48 89 c7                  mov     %rax,%rdi
11df:      b8 00 00 00 00             mov     $0x0,%eax
11e4:      e8 97 fe ff ff             call    1080 <printf@plt>
11e9:      8b 55 ec                  mov     -0x14(%rbp),%edx
11ec:      48 8d 45 ec                lea     -0x14(%rbp),%rax
11f0:      48 89 c6                  mov     %rax,%rsi
11f3:      48 8d 05 0e 0e 00 00      lea     0xe0e(%rip),%rax      # 2008
11fa:      48 89 c7                  mov     %rax,%rdi
11fd:      b8 00 00 00 00             mov     $0x0,%eax
1202:      e8 79 fe ff ff             call    1080 <printf@plt>
1207:      48 8b 45 f0                mov     -0x10(%rbp),%rax
120b:      8b 08                      mov     (%rax),%ecx
120d:      48 8b 55 f0                mov     -0x10(%rbp),%rdx
1211:      48 8d 45 f0                lea     -0x10(%rbp),%rax
1215:      48 89 c6                  mov     %rax,%rsi
1218:      48 8d 05 01 0e 00 00      lea     0xe01(%rip),%rax      # 2020
121f:      48 89 c7                  mov     %rax,%rdi
1222:      b8 00 00 00 00             mov     $0x0,%eax
1227:      e8 54 fe ff ff             call    1080 <printf@plt>
122c:      90                          nop
122d:      48 8b 45 f8                mov     -0x8(%rbp),%rax
1231:      64 48 2b 04 25 28 00      sub     %fs:0x28,%rax
1238:      00 00
123a:      74 05                      je      1241 <main+0xb8>
123c:      e8 2f fe ff ff             call    1070 <__stack_chk_fail@plt>
1241:      c9                          leave
1242:      c3                          ret

```

Disassembly of section .fini:

```

00000000000001244 <_fini>:
    1244:      f3 0f 1e fa      endbr64
    1248:      48 83 ec 08      sub     $0x8,%rsp
    124c:      48 83 c4 08      add     $0x8,%rsp
    1250:      c3              ret

```

Function at start address: `_start`

The start function executes **before** the main function, and set up the environment for the program to run. Such as setting up the stack, heap and other components. It also calls the main function, and when the main function returns, it handles the program's exit.

4. Run the program several times and see the address change.

The reason the address changes each time is because the address is randomized for each execution. This is simply a security feature, to make it harder for a bad actor to predict where certain parts of the program will be stored. Also known as ASLR (Address Space Layout Randomization).

Task 4: The stack.

1. Compile the code with gcc.

Disclaimer: Considering im using the newest version of ubuntu as operating system, some changes was made to the code to make it compile without errors. Using `@ %p`" instead of `@ 0x%08x` i do belive it should print the same result in this case.

2. Determine default size of stack for your Linux system.

```

$ ulimit -s
8192

```

Default stack size: 8192 KB (8 MB)

3. Results from running the code:**

```

...
func() frame address @ 0x7ffded2277d0
func() frame address @ 0x7ffded2277b0

```

```
func() frame address @ 0x7ffded227790
Segmentation fault (core dumped)
```

This fault is simply the program **exceeding the stack limit**, by calling the function recursively without a base case to stop it.

4. What does `grep func | wc -l` catch?

This line `grep func` grabs lines containing the word *func*, and `wc -l` is a command that counts the number of lines.

So in this case we got:

```
261712
```

meaning that the function `func` was called 261712 times before the stack overflowed and caused a segmentation fault.

5. How much stack memory *bytes* does each recursive call use?

To calculate bytes of a stack we can simply take the difference between two consecutive frame addresses. It is important to note that in stack we start from high adress to low adress. So we must subtract the lower adress from the higher adress. Which is uncommin in a mathematical context.

```
func() frame address @ 0x7ffe0a822920
func() frame address @ 0x7ffe0a822900
```

if we take $0x7ffe0a822920 - 0x7ffe0a822900 = 0x20$

So each recursive call uses `0x20` bytes of memory, which is `32` bytes in decimal. This can easily be calculated by using a hex to decimal converter. In this case I used rapidtables

From	To
Hexadecimal ▾	Decimal ▾

Enter hex number

0x20	16
------	----

= Convert **× Reset** **↕ Swap**

Decimal number (2 digits)

32	10
----	----